

Foundations of C++


Bjarne Stroustrup

Texas A&M University



Overview

- Memory and objects
 - Construction and destruction
- Containers
 - Copy and Move
- Resources and RAII
- Class hierarchies
- Algorithms
- Compile-time computation
 - Type functions
- Concurrency
- Not: Casts, macros, pointer arithmetic, how to write bad code



Foundations of C++


- This is a talk about programming techniques
 - And language support for such techniques
- I present an industrial programmer's view of C++
 - No Γρεεκ Λεττερς
 - No Grammar
- I present fundamental examples
 - Not language details
 - Not legacy techniques/code
 - There are hundreds of millions of lines of code relying on the techniques I mention
 - An idealistic view: progress is necessary and possible
- I don't focus on the new C++ features
 - C++ is not (just) a series of historical strata

Complexity

- C++ is huge
 - But so are other language used for production code
- Complexity goes somewhere
 - Language, library, application, infrastructure
- The very notion of programming is changing/fracturing
 - Library users
 - Scripters
 - System builders
 - Infrastructure builders
 - Embedded systems builders
 - ...

*proxies for size comparisons:
spec #words
library #types (non-'plumbing')*

Portable C++



Herb Sutter

Scale: C++ language: 400 pages; C++ standard library: 750 pages

Portable C++

language

library

2008 .NET FX + VS Pro Libs

Java SE 7

2008 .NET FX (only)

Java 7 (2011)

C# 3.0 (2008)

C++11

C++11

Herb Sutter


C++

- ISO/IEC 14882-2011 aka C++11, formerly “C++0x”
- Basics:
 - A simple and direct mapping to hardware
 - Zero-overhead abstraction mechanisms
- Supports
 - Classical systems programming
 - Infrastructure applications
 - resource-constrained and mission-critical
 - Light-weight abstraction
 - A type-rich style of programming
 - C++ supports type-safe programming with a non-trivial set of types.
 - And more

Other concerns

- Most of what is important to software development organizations don't show in code fragments
 - Tool chains
 - Stability ***and*** progress
 - Interoperability with other languages
 - Availability of libraries
 - Availability of trained developers


Memory model



Memory is sequences of objects addressed by pointers

Memory model (built-in type)

- char
- short
- int
- long
- (long long)
- float
- double
- long double
- T* (pointer)
- T& (implemented as pointer)



Memory model (“ordinary” class)

```
class Point {  
    int x, y;  
    // ...  
};
```

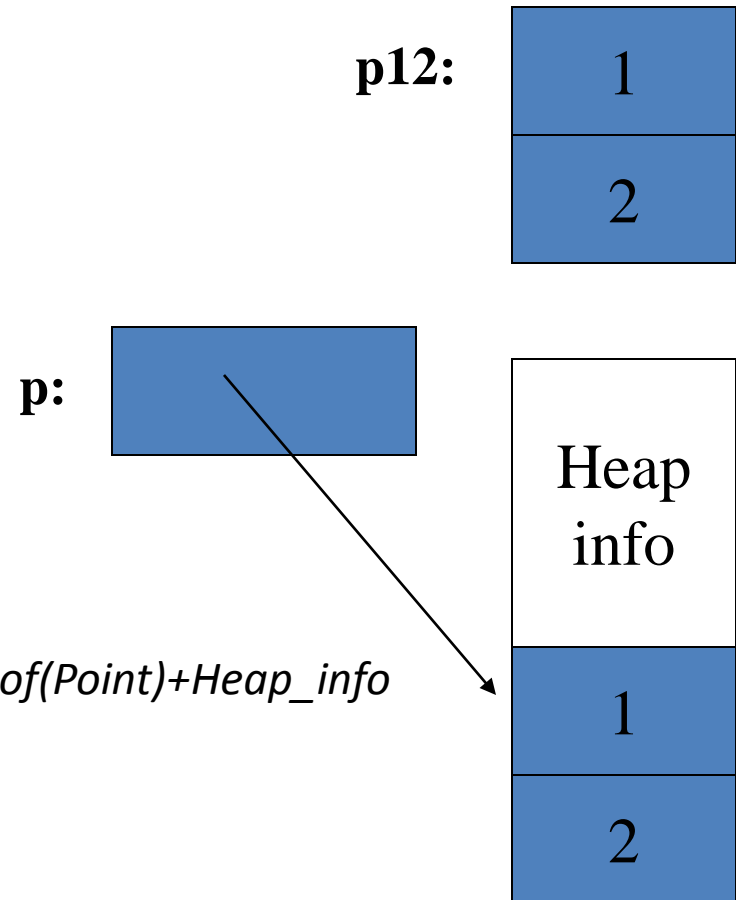
*// sizeof(Point)==2*sizeof(int)*

```
Point p12 {1,2};
```

```
Point* p = new Point{1,2};
```

// memory used for “p”:sizeof(Point)+sizeof(Point)+Heap_info*

- Simple Composition




Memory model – class hierarchy

```
class B {  
    int b;  
};
```



```
class D : public B {  
    int d;  
};
```



```
B x;  
D y;
```


- Simple composition

Memory model (polymorphic type)

```
class Shape {
public:
    virtual void draw() = 0;
    virtual Point center() = 0;
    // ...
};
```

Shape* p = new Circle{{x,y},r};




p:



Heap
info

vptr


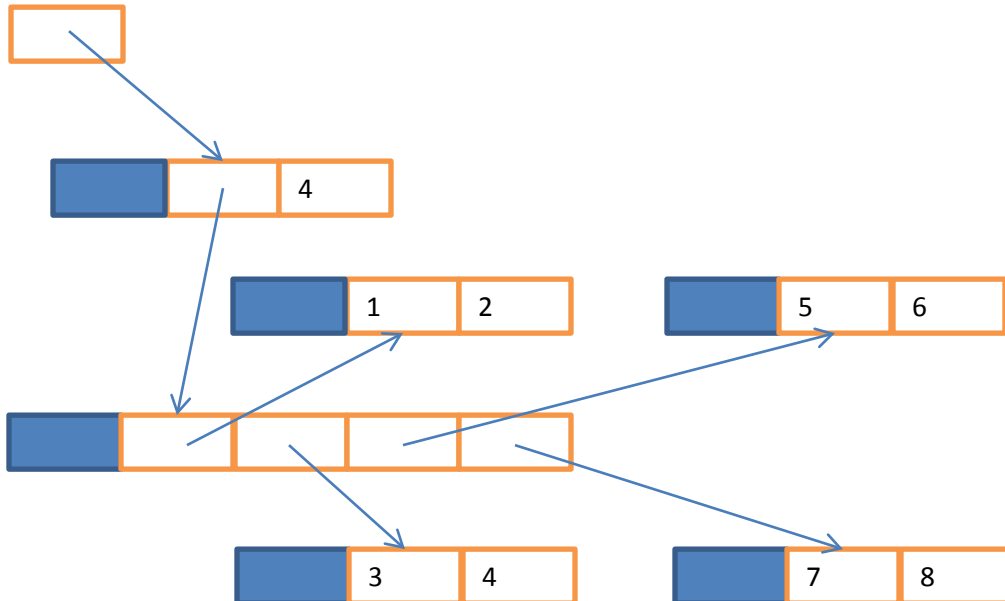
vtbl:



Use compact layout

- `vector<Point> vp = { Point{1,2}, Point{3,4}, Point{5,6}, Point{7,8} };`

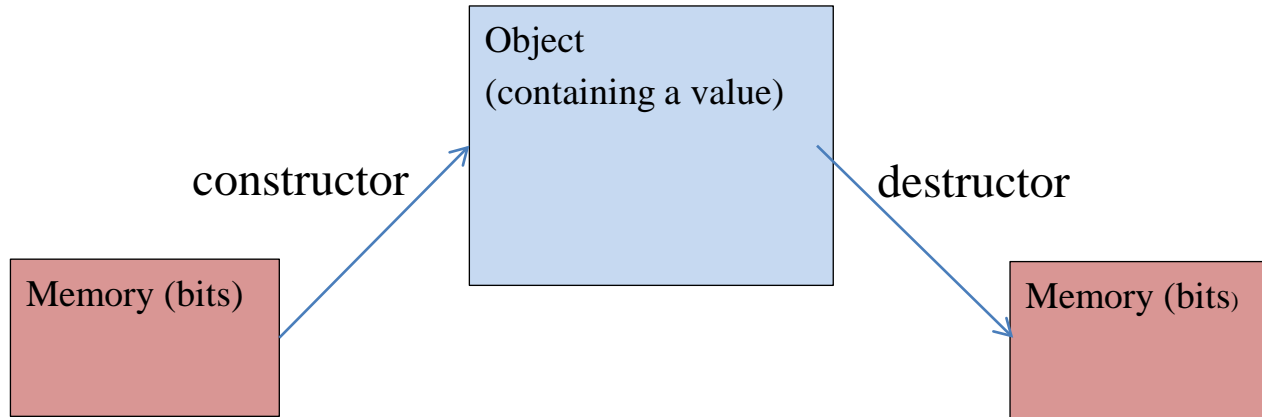
“True OO” style:



A loss


- Many students and developers don't understand the language-to-machine mapping
 - To them, it's "magic"
 - In the context of infrastructure projects, that's a significant problem

Constructors and destructors



- Constructor: make an object from memory
 - An object holds a value
 - Has a type
 - Has an interface
 - Has meaning
- Destructor: make an object (back) into memory
 - Memory is just interpreted bits

A resource handle



- Examples
 - Containers: vector, list, map, ...
 - Smart pointers: `unique_ptr`, `shared_ptr`, `delayed_value`, `remote_object`, ...
 - Locks, thread handles, sockets, iostreams
 - File handle
 - ...

Vector: the archetypical resource handle

- Slight simplification of `std::vector`

```

template<typename T>                                // T is the element type
class Vector {
public:
    Vector();                                         // default constructor: make empty vector
    Vector(int n);                                   // constructor: initialize to n elements
    Vector(initializer_list<T>);                   // constructor: initialize with element list
    ~Vector();                                       // destructor: deallocate elements
    int size();                                     // number of elements
    T& operator[](int i);                           // access the ith element
    void push_back(const T& x);                     // add x as a new element at the end
    T* begin();                                     // fist element
    T* end();                                       // one-beyond-last element
private:
    int sz;                                         // number of elements
    T* elem;                                       // pointer to sz elements of type T
};

```

Vector use

- Simple use of vectors

```
void f(Vector<string>& vs)
{
    Vector<int> sizes;                // empty vector: sizes.size()==0
    for (auto x : vs)                // loop through all elements of vs
        sizes.push_back(x.size());  // add element to vector (grow)

    if (0<vs.size())                 // check size
        vs[0] = "Whatever!";        // subscripting

    // ...
}
```

Vector use

- Simple use of vectors

```
int main()
{
    f({"Wheeler", "Wilkes", "Radcliffe", "Appleton", "Rutherford"});
    Vector<string> places(10); // 10 empty strings
    places[2] = "Cambridge";
    // ...
    f(places);
}
```

Constructor

- N default elements

```
template<typename T>
Vector<T>::Vector(int n) // make a vector with n elements of default value
    :sz{n},
    elem{allocate<T>(sz)}    // allocate space for sz elements of type T
{
    if (sz<0) throw std::runtime_error{"negative Vector size"};
    std::uninitialized_fill(elem,elem+sz,T{});    // initialize to default
}
```

- The **allocate<T>()** function is a simplification of the standard-library **allocator** mechanism to make the examples fit on slides

Destructor

- Essential:
 - release resource (in this case free memory)

```
template<typename T>
```

```
Vector<T>::~~Vector()           // destructor releases resources acquired
```

```
{
```

```
    destroy<T>(elem,n); // invoke member destructors, then deallocate elem[]
```

```
}
```

- Note: the elements typically have destructors
 - E.g. **Vector<Vector<string>>**
 - Explicit invocation of destructors is extremely rare
 - basically only in sophisticated container implementations

Initializer-list constructor

- A class can have many constructors

```
template<typename T>
Vector<T>::Vector(std::initializer_list<T> lst)    // elements from the list
    :sz{lst.size()},
    elem{allocate<T>(sz)}
{
    std::uninitialized_copy(lst.begin(), lst.end(), elem);
}
```

A matched blend of techniques

- The standard library containers (e.g., **vector**, **map**, **set**, and **list**) :
 - classes for separating interfaces from implementations
 - constructors for establishing invariants, including acquiring resources
 - destructors for releasing resources
 - templates for parameterizing types and algorithms with types
 - mapping of source language features to user-defined code
 - e.g. `[]` for subscripting, the **for**-loop, **new/delete** for construction/destruction on the free store, and the `{}` lists.
 - use of half-open sequences, e.g. `[begin():end())`, to define for-loops and general algorithms.
 - Use of standard-library facilities to simplify specification and implementation
- This abstraction from “memory” to “containers of objects” carries no overheads
 - beyond the code necessarily executed for memory management, initialization, and error checking.

Absolutely minimal overheads

- There is no data stored in a **Vector** object
 - beyond the two named members (three in **std::vector**)
- The element type need not be part of a hierarchy
 - the only requirements on a template argument are imposed by its use
 - “duck typing.”
- **Vector** operations are not dynamically resolved
 - Not **virtual**
 - Simple operations, such as **size()** and **[],** are typically inlined
- A Vector is allocated where needed
 - On stack, in objects
- A vector is accessed directly
 - Not through a handle (*it is a handle*)

Copy and move

- Copy constructor

```
Vector capitals {"Helsinki", "København", "Riga", "Tallinn"};  
Vector c2 = capitals;    // error: no copy defined for Vector  
// By default, you can copy only objects with “simple representations.
```

// So we define a suitable copy:

```
template<typename T>  
Vector<T>::Vector(const Vector& v)    // copy constructor  
    : sz{v.sz}, elem{allocate<T>(v.sz)}  
{  
    std::uninitialized_copy(v.begin(),v.end(), elem);  
}
```

Copy and move

- Copy assignment

```
template<typename T>
Vector<T>& Vector<T>::operator=(Vector<T>& v)  // copy assignment
{
    Vector<T> tmp {v};                        // copy v
    destroy<T>(elem,sz);                       // delete old elements
    elem = tmp.elem;                           // "steal" tmp's representation
    sz = tmp.sz;
    tmp.elem = nullptr;                        // tmp will be empty when destroyed
    tmp.sz = 0;
    return *this;
}
```

Copy and move

```
Vector<int*> find_all(Vector<int>& v, int val) // find all occurrences of val in v
{
    Vector<int*> res;
    for (int& x : v)
        if (x==val)
            res.push_back(&x);    // add the address of the element to res
    return res;
}
```

```
void test()
{
    Vector<int> lst {1,2,3,1,2,3,4,1,2,3,4,5};
    for (int* p : find_all(lst,3))
        cout << "address: " << p << ", value: " << *p << "\n";
    // ...
}
```

Copy and move

- Move constructor

```
template<typename T>
```

```
Vector<T>::Vector(const Vector&& v)           // move constructor
```


```
    : sz{v.sz}, elem{v.elem}           // grab v's elements
```

```
{
```

```
    v.elem = nullptr;                   // make v empty
```

```
    v.sz = 0;
```

```
}
```



Copy and move

- Move assignment

```
template<typename T>
```

```
Vector<T>& Vector<T>::operator=(Vector<T>&& v)  // move assignment
```

```
{
```

```
    destroy<T>(elem,sz);           // delete old elements
```

```
    elem = v.elem;                 // grab v's elements
```


```
    sz = v.sz;
```

```
    v.elem = nullptr;              // make v empty
```

```
    v.sz = 0;
```

```
    return *this;
```

```
}
```



Vector

- Copy and move declarations added to Vector

```
template<typename T>           // T is the element type
class Vector {
public:
    // ...
    Vector(const Vector&);      // copy constructor
    Vector(Vector&&);           // move constructor
    Vector& operator=(const Vector&); // copy assignment
    Vector& operator=(Vector&&);   // move assignment
    // ...
};
```

Vector

- And of course, a user doesn't have to implement **Vector**
 `#include<vector>`
will get an even more flexible and efficient version
- BUT: all the techniques and language facilities are available for all to use for their own abstractions

Resources and errors

- Exceptions
- Resources
- RAII



Error Handling: Exceptions

```
void do_task(int i)
{
    if (i==0) throw std::runtime_error{"do_task() of zero"};
    if (i<0) throw Bad_arg{i};
    // do the task and return normally
}

void task_master(int i)
{
    try {
        do_task(i);
        // ...
    }
    catch (Bad_arg a) {
        cout << "do_task() of negative" << a.val << "\n";
    }
}
```

Real-world constraints

- Sometimes, you can't use exceptions
 - Hard real time
 - Messy old code
- Implications
 - Duplication of styles and components
 - Complexity
 - Confusion
- This happens again and again
 - for different programming techniques
 - For different language features
 - A major source of complexity

Resources and Errors

- Many (most?) resources are not just memory
 - A non-memory resource requires a release operation
 - Not just freeing of memory

// unsafe, naïve use:

```
void f(const char* p)  
{  
    FILE* f = fopen(p,"r");           // acquire  
    // use f  
    fclose(f);                       // release  
}
```

Resources and Errors

```
// naïve fix:

void f(const char* p)
{
    FILE* f = 0;
    try {
        f = fopen(p, "r");
        // use f
    }
    catch (...) {           // handle every exception
        if (f) fclose(f);
        throw;
    }
    if (f) fclose(f);
}
```

RAII (Resource Acquisition Is Initialization)

```
// use an object to represent a resource
class File_handle {    // belongs in some support library
    FILE* p;
public:
    File_handle(const char* pp, const char* r)
        { p = fopen(pp,r); if (p==0) throw File_error(pp,r); }
    File_handle(const string& s, const char* r)
        { p = fopen(s.c_str(),r); if (p==0) throw File_error(s,r); }
    ~File_handle() { fclose(p); }    // destructor
    // copy operations
    // access functions
};

void f(string s)
{
    File_handle fh {s, "r"};
    // use fh
}
```

RAII

- For all resources
 - Memory (done by **std::string**, **std::vector**, **std::map**, ...)
 - Locks (e.g. **std::unique_lock**), files (e.g. **std::fstream**), sockets, threads (e.g. **std::thread**), ...

```
std::mutex m;           // a resource  
int sh;                // shared data
```

```
void f()  
{  
    // ...  
    std::unique_lock<mutex> lck {m};    // grab (acquire) the mutex  
    sh+=1;        // manipulate shared data  
}                                     // implicitly release the mutex
```

Simplify control structure

- Prefer algorithms to unstructured code



Algorithms

- Messy code is a major source of errors and inefficiencies
- We must use more “standard” well-designed and tested algorithms
- The C++ standard-library algorithms are expressed in terms of half-open sequences [**first:last**)
 - For generality and efficiency

```
void f(vector<int>& v, list<string>& lst)
```

```
{
```

```
    sort(v.begin(),v.end());
```

```
    auto p = find(lst.begin(),lst.end(),"Aarhus"); // find "Aarhus" in lst:
```

```
    if (p!=lst.end()) { // found: *p=="Aarhus"
```

```
        // ...
```

```
    else // not found *p!="Aarhus"
```

```
        // ...
```

```
}
```

Algorithms

- Simple, efficient, and general implementation
 - For any forward iterator
 - For any (matching) value type

```
template<typename Iter, typename Value>  
Iter find(Iter first, Iter last, Value val)  
    // find first p in [first:last) so that *p==val  
{  
    while (first!=last && *first!=val)  
        ++first;  
    return first;  
}
```

Algorithms

- Parameterization with criteria, actions, and algorithms
 - Essential for flexibility and performance

```
void g(vector< string>& vs)
{
    auto p = std::find_if(vs.begin(), vs.end(), Less_than{"Griffin"});

    if (p!=vs.end()) {    // found: *p<"Griffin"
        // ...
    }
    else {                // not found *p>="Griffin"
        // ...
    }
    // ...
}
```

Algorithms

- The implementation is still trivial

```
template<typename Iter, typename Value>
Iter find_if(Iter first, Iter last, Predicate pred)
    // find first p in [first:last) so that pred(*p)
{
    while (first!=last && !pred(*first))
        ++first;
    return first;
}
```

Algorithms: function objects

- General function object
 - Can carry state
 - Easily inlined

```
struct Less_than {  
    String s;  
    Less_than(const string& ss) :s{ss} {} // the value to compare against  
    bool operator(const string& v) const { return v<s; } // the comparison  
};
```

Lambda notation

- We can let the language write the function object

```
auto p = std::find_if(vs.begin(),vs.end(),  
                    [](const string& v) { return v<"Griffin"; } );
```


Container algorithms

- The C++ standard-library algorithms are expressed in terms of half-open sequences [first:last)
 - For generality and efficiency
 - If you find that verbose define container algorithms

```
namespace Extended_STL {  
    template<typename C>  
    void sort(C& c) { std::sort(c.begin(),c.end()); }  
    // ...  
}
```

Compile-time Computation

- Type-rich computation at compile time.
 - *Efficiency*: To pre-calculate a value (often a size).
 - Simple cases (only) done by an optimizer.
 - *Type-safety*: To compute a type at compile time.
 - *Simplify concurrency*: you can't have a race condition on a constant.
- No just error-prone macro hacking



Type-rich compile-time computation

- Just like other code
 - Except it is executed by the compiler
 - To do anything interesting we need a type system

```
struct City { double x, y };
```

```
constexpr double csqrt(double) { /* calculate square root */ }
```

```
constexpr double square(double d) { return d*d; }
```

```
constexpr double dist(City c1, City c2)
```

```
{
```

```
    return csqrt(square(abs(c1.x-c2.x))+square(abs(c1.y-c2.y)));
```

```
}
```


```
constexpr double d = dist(NewYork,Boston);    // a simple use
```


Unit checking: SI Units

- Units are effective and simple:

```
Speed sp1 = 100m/9.8s;           // very fast for a human
Speed sp2 = 100m/9.8s2;          // error (m/s2 is acceleration)
Speed sp3 = 100/9.8s;            // error (speed is m/s and 100 has no unit)
Acceleration acc = sp1/0.5s;     // too fast for a human
```

- and essentially free (in C++11)
 - Compile-time only
 - No run-time overheads



Type Functions

- The standard library depends on
 - Compile-time selection of types
 - Compile-time calculation of values
 - Compile-time selection of algorithms
- So does much other code
- Sometimes, referred to as “Template Meta-programming”
 - Keep it simple
- The key notion is a “Type function”
 - takes at least one type argument or returns at least one type
 - **sizeof(T)**
 - **SameType<T,U>**
 - **Value_type<Iter>**

Type Functions

- Functions that answer questions about types or return types

```
template<typename Cont>
```

```
void sort(Cont& c)           // sort container of type Cont
```

```
{
```

```
    if (Has_random_access<Cont>())    // ask about Cont's properties
```

```
        sort(c.begin(),c.end());
```

```
    else {
```

```
        vector<Value_type<Cont>> v {c.size()}; // get an associated type from Cont
```

```
        copy(c.begin(),c.end(),v);           // copy, sort, and copy back
```

```
        sort(v);
```

```
        copy(v.begin(),v.end(),c);
```

```
    }
```

```
}
```

Type Functions

- I really wanted to overload on “concepts” (compiler-supported predicates on sets of types and values):
 template<Random_access_container> ...
 template<Bidirectional_access_container> ...
 template<Forward_access_container> ...
- But for now I will show how to use type functions
 - As widely used in current C++
 - How to non-intrusively add properties to types
 - Using “traits”

Type Functions

- Traits classes (an important technique/workaround)
 - A general mechanism for adding non-intrusively properties to types

```
template<typename Cont>
```

```
struct container_traits {           // general/default traits  
    using value_type = typename Cont::value_type;      // member type  
    using access_category = typename Cont::access_tag; // member type  
    // ...  
};
```

- Why not always use **Cont::value_type**?
 - Because **T[N]::value_type** is invalid syntax

Type Functions

- We use a template alias to return a type:

```
template<typename Cont>  
using Value_type = typename container_traits<Cont>::value_type;
```

```
template<typename Cont>  
using Access_category =  
    typename container_traits<Cont>::access_category;
```

- So, `Value_type<Vector<double>>` is `double`

Type Functions

- We use `constexpr` function templates to return values

```
template <typename T, typename U>
constexpr bool Is_same()
{
    return is_same<T, U>::value;  // is_same is a std type trait (an intrinsic)
}
```

```
template<typename Cont>
constexpr bool Has_random_access<Cont>()
{
    return Is_same<Access_category<Cont>, random_access_tag>();
}
```

- `Has_random_access<Vector<double>>` is true

Type Functions

- Functions that answer questions about types or return types

```
template<typename Cont>
```

```
void sort(Cont& c)           // sort container of type Cont
```

```
{
```

```
    if (Has_random_access<Cont>())    // ask about Cont's properties
```

```
        sort(c.begin(),c.end());
```

```
    else {
```

```
        vector<Value_type<Cont>> v {c.size()}; // get an associated type from Cont
```

```
        copy(c.begin(),c.end(),v);           // copy, sort, and copy back
```

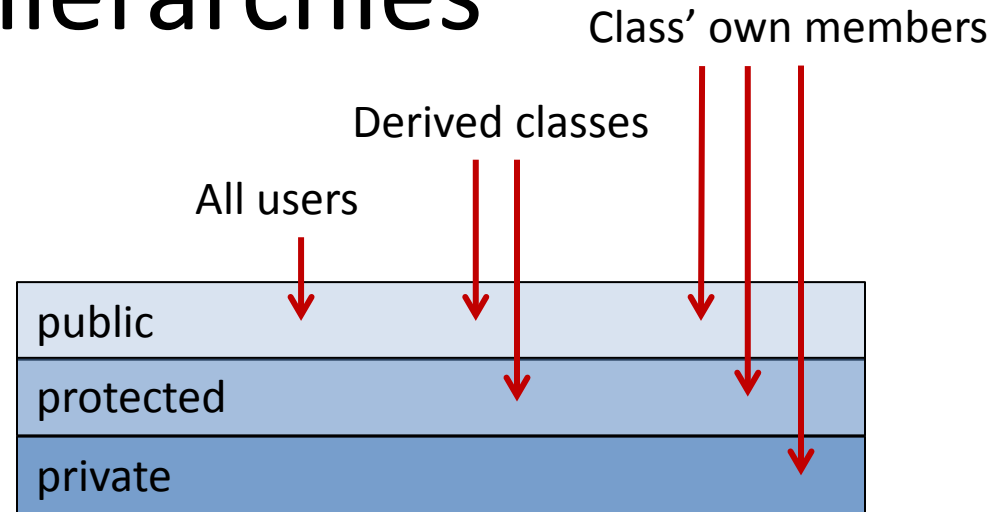
```
        sort(v);
```

```
        copy(v.begin(),v.end(),c);
```

```
    }
```

```
}
```


Class hierarchies



- Protection model
- No universal base class
 - an unnecessary implementation-oriented artifact
 - imposes avoidable space and time overheads.
 - encourages underspecified (overly general) interfaces
- Multiple inheritance
 - Interface and implementation
 - Abstract classes provide the most stable interfaces
- Minimal run-time type identification
 - `dynamic_cast<D*>(pb)`
 - `typeid(p)`

“Paradigms”

- Much of the distinction between object-oriented programming and generic programming is an illusion
 - based on a focus on language features
 - incomplete support for a synthesis of techniques
 - The distinction does harm
 - by limiting programmers, forcing workarounds

```
template<typename Cont>  
void draw_all(Cont& c)  
{  
    for_each(c.begin(),c.end(), [](Shape* p) { p->draw(); } )  
}
```

Concurrency

- There are many kinds
- Stay high-level
- Stay type-rich



Type-Safe Concurrency

- Programming concurrent systems is hard
 - We need all the help we can get
 - C++11 offers type-safe programming at the threads-and-locks level
 - Type safety is hugely important
- threads-and-locks
 - is an unfortunately low level of abstraction
 - is necessary for current systems programming
 - That's what the operating systems offer
 - presents an abstraction of the hardware to the programmer
 - can be the basis of other concurrency abstractions

Threads

```
void f(vector<double>&);           // function

struct F {                        // function object
    vector<double>& v;
    F(vector<double>& vv) :v{vv} { }
    void operator()();
};

void code(vector<double>& vec1, vector<double>& vec2)
{
    std::thread t1 {f,vec1};       // run f(vec1) on a separate thread
    std::thread t2 {F{vec2}};     // run F{vec2}() on a separate thread
    t1.join();
    t2.join();
    // use vec1 and vec2
}
```

Thread – pass argument and result

```
double* f(const vector<double>& v);  // read from v return result
double* g(const vector<double>& v);  // read from v return result

void user(const vector<double>& some_vec)      // note: const
{
    double res1, res2;
    thread t1 { [&]{ res1 = f(some_vec); }};  // lambda: leave result in res1
    thread t2 { [&]{ res2 = g(some_vec); }};  // lambda: leave result in res2
    // ...
    t1.join();
    t2.join();
    cout << res1 << ' ' << res2 << '\n';
}
```

async() – pass argument and return result

```
double* f(const vector<double>& v);  // read from v return result
```

```
double* g(const vector<double>& v);  // read from v return result
```

```
void user(const vector<double>& some_vec)      // note: const  
{  
    auto res1 = async(f,some_vec);  
    auto res2 = async(g,some_vec);  
    // ...  
    cout << *res1.get() << ' ' << *res2.get() << '\n';  // futures  
}
```

- Much more elegant than the explicit thread version
 - And most often faster

No garbage collection needed

- Apply these techniques in order:
 1. Store data in containers
 - The semantics of the fundamental abstraction is reflected in the interface
 - Including lifetime
 2. Manage all resources with resource handles
 - RAI
 - Note: non-memory resources
 3. Use “smart pointers”
 - They are still pointers
 4. Plug in a garbage collector
 - For “litter collection”
 - C++11 specifies an interface
 - Can still leak non-memory resources

Type safety

- C++ is not guaranteed to be statically type safe
 - “C is a strongly typed; weakly checked, language” – DMR
- A language designed for general and performance critical systems programming with the ability to manipulate hardware cannot be.
- Problems
 - untagged unions
 - explicit type conversions (casts)
 - arrays without (guaranteed) range checks
 - ability to deallocate a free store (heap) object while holding on to a pointer allowing for post-allocation access.
 - ability to deallocate an object not allocated on the free store

Challenges



- Obviously, C++ is not perfect
 - How can we make programmers prefer modern C++ styles over low-level (C-style) code, which is far more error-prone and harder to maintain, yet no more efficient?
 - How can we make C++ a better language given the Draconian constraints of C and C++ compatibility?
 - How can we improve and complete the techniques and models (incompletely and imperfectly) embodied in C++?
- In the context of C++, solutions that eliminate major C++ strengths are not acceptable
 - Compatibility (link, source code)
 - Performance
 - Portability
 - Range of application areas

Challenges

- Close more type loopholes
 - in particular, find a way to prevent misuses of **delete** without spoiling RAI
- Simplify concurrent programming
 - in particular, provide some higher-level concurrency models as libraries
- Simplify generic programming
 - in particular, introduce simple and effective concepts
- Simplify programming using class hierarchies
 - in particular, eliminate use of the visitor pattern
- Better support for combinations of object-oriented and generic programming
- Make exceptions usable for hard-real-time projects
 - that will most likely be a tool rather than a language change
- Find a good way of using multiple address spaces
 - as needed for distributed computing
 - would probably involve defining a more general module mechanism that would also address dynamic linking, and more.
- Provide many more domain-specific libraries
- Develop a more precise and formal specification of C++

Questions?

C++: A light-weight abstraction
programming language



Key strengths:

- software infrastructure
- resource-constrained applications

Practice type-rich
programming